



NoSQL Document Database



Getting Started | Creating CRUD



Ricardo de Luna Galdino
Software Engineer

github.com/ricardogaldino/microlearning-ravendb

Agile learning culture

Microlearning:

- It is a teaching methodology that subdivides a subject into **smaller doses of content**, with quick activities, helping to understand and retain this content;

Knowledge Pills:




- These are **small contents** presented to the professional so that he **can assimilate them in a more focused and objective** way, improving efficiency and enhancing the results obtained;
- **Microlearning is composed** of several **Knowledge Pills**;

⇒ 01

Introduction

Introduction to the Database
Non-Relational and **NoSQL**



A brief history...

- The term NoSQL **appeared** in mid-**1998** by Software Engineer “**Carlo Strozzi**” and initially would be “**NoREL**”, because he was developing a database that would not be relational;
- “**No SQL**”, for some it became “**Not Only SQL**” because it does not replace, in all aspects, relational databases, in fact, currently (2023) they complement each other;
- It was driven by **Google** (BigTable, 2004) and **Amazon** (Dynamo, 2007) when they **sought a new form of data storage** that **could scale horizontally**, with the creation of clusters. It was a **good idea**, both from a **technical** point of view and from an **economic** point of view;



Carlo Strozzi

Concepts about NoSQL

Non-relational

- The term “**Non-Relational**” refers to DBs that do not follow the “Relational” architecture, that is, they do **not use** the **Tables**, **Rows** and **Columns** schema when storing data;
- Data can be stored, for example: as **Key/Value** pairs, as **Documents** (XML, JSON, ..) or as **Graphs** (nodes and edges);

NoSQL

- Originally the term “NoSQL” referred to databases that do **not use SQL** (Structured Query Language) **for queries**;

References...

- Microsoft, 2023. “Non-relational data and NoSQL”.
- Microsoft, 2023. “What are NoSQL databases?”.
- Amazon, 2023. “What is NoSQL?”.

Advantages of NoSQL

- **Flexibility:** they can store and combine any type of data, whether structured or unstructured, unlike relational databases, which only store in a structured way;
- **Scalability:** they support growth in volume of data and simultaneous access well because they are built to scale horizontally (distributed environments (clusters)), unlike relational databases, which generally scale vertically;
- **Availability:** they are efficient in data replication, that is, if one or more servers go down, another is able to continue the work. With this, they avoid the loss of data and, consequently, of dissatisfied customers;
- **High performance:** they are built to have great performance, measured by the transfer rate (amount of data (bits) transferred per time (seconds)) and latency (time required to make a request and receive the response);
- **Dynamic schema:** accepts changes in the data structure (schema) without major problems;
- **Open source:** Because they are open source, they require no licensing fees and can be deployed cost-effectively;

Disadvantages of NoSQL

- **Data consistency** (integrity constraints to ensure data accuracy and reliability) is not one of your priorities. This responsibility will be up to the application that will use NoSQL;

For example, it natively allows data duplication;

- **Data normalization** is also not one of your priorities. It uses denormalized data models (it contains all the necessary information even if it is repeated);
- They were not made for **complex queries** with reference between several “**joins**” entities;

- Doesn't contain as many advanced features and tools as relational databases;
- There is a lack of **standardization** and incompatibility with each other. Each uses its query language and implements concepts in its own way;

NOTE!

Relational Databases prioritize these items natively and that's why they are still necessary and work together with NoSQL databases!

NewSQL movement! Will it be the future?

- **NewSQL** databases seek to promote the same improvement in performance and scalability as NoSQL systems, without giving up the benefits of relational databases and the SQL language;
- The term was first used by Data Engineer “**Matthew Aslett**” in a **2011** research paper discussing the emergence of a new generation of database management systems;
- The “BOOM” has not yet occurred, but what we see is, little by little, the **inclusion** of **Relational** Database (SQL) **resources in NoSQL** Databases and the same in reverse;




Matthew Aslett

⇒ 02

RavenDB

Concepts



RavenDB - Concepts



Oren Eini

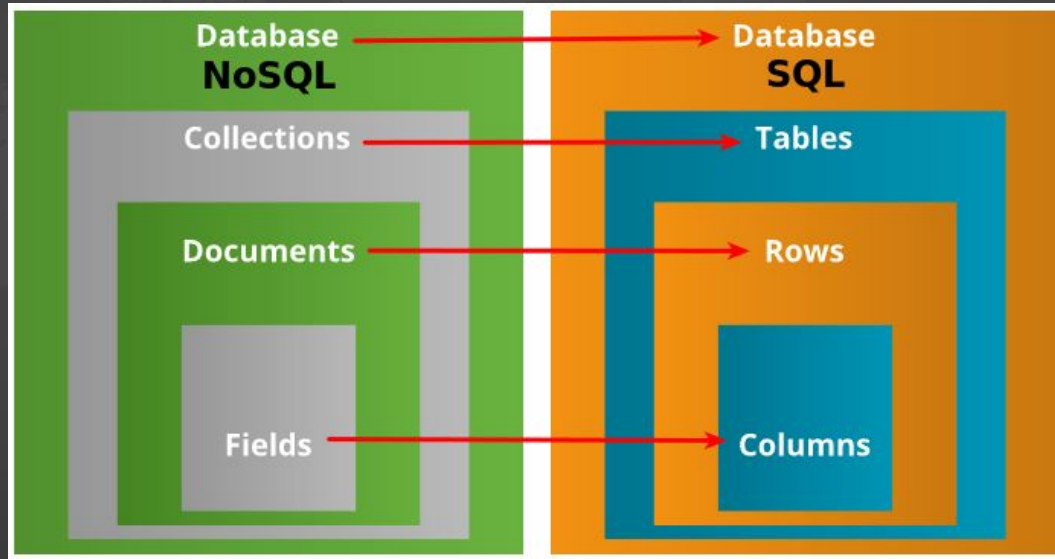
- It is a **document**-oriented **NoSQL** database (default **JSON**), open-source, written in C# and with support for ACID transactions;
- Developed in **2010**, under the leadership of **Oren Eini**, by **Hibernating Rhinos Ltd**;
- **Multi-platform**, running on Windows, Linux, MacOS, Docker and Raspberry Pi;
- In **2019** it began to be offered, as a **Cloud Service**, on AWS, Azure and GCP;
- Uses the **RQL** (Raven Query Language) **for queries**;
- **LINQ** (Language Integrated Query) support in **C#**;
- Has **clients** in C#, C++, Java, Node.js, Python, Ruby, PHP, Elixir (Erlang) and Go;
- Do **dynamic indexing**. All queries use indexes. If it does not exist, an index is **automatically created**;

References...

- RavenDB 2023. "Documentation".
- RavenDB 2023. "RDL (Raven Query Language)".

RavenDB - Architecture



We can **compare** the architecture of a “**NoSQL**” **Document** DB with a “**SQL**” Relational DB:



RavenDB - Modeling

A DB of the “**Document**” type is oriented to **aggregates** (a term that comes from **Domain-Driven Design**), which means, in a simplified way, that each document is a collection of related objects, treated as a single unit over a context.

Ex: Invoice, its items and other data would be a **single denormalized "JSON" document**.

RECEBEMOS DE: OS PRODUTOS CONSTANTES DA NOTA FISCAL INDICADA AO LADO		NF-e Nº 000.000.000 SÉRIE 000	
DATA DE RECEBIMENTO	IDENTIFICAÇÃO E ASSINATURA DO RECEBEDOR		
<div>IDENTIFICAÇÃO DO EMITENTE</div> <div> 0 - - Telefone: () -</div>		<div>DANFE</div> <div>Documento Auxiliar da Nota Fiscal Eletrônica 0 - Entrada 1 - Saída Nº 000.000.000 Série 000 FL 1 / 1</div> <div></div> <div>CHAVE DE ACESSO 0098.1200.0000.0000.0010.0000.0009</div> <div>Consulta de autenticidade no portal nacional da NF-e www.nfe.fazenda.gov.br/portal ou no site da Sefaz Autorizadora</div>	
NATUREZA DA OPERAÇÃO		DADOS DA NF-e	
INSCRIÇÃO ESTADUAL	INSC. ESTADUAL SUBST. TRIBUTÁRIO	CNPJ	
DESTINATÁRIO / REMETENTE		CNPJCPF	DATA DA EMISSÃO
NOME RAZÃO SOCIAL			30/12/1999
ENDEREÇO	BAIRRO/DISTRITO	CEP	DATA DA ENTRADA/SADA
MUNICÍPIO	FONE/FAX	UF	HORA DE SAÍDA
INSCRIÇÃO ESTADUAL			
FATURA			
CÁLCULO DO IMPOSTO			
BASE DE CÁLCULO DE ICMS		VALOR DO ICMS	
0,00		0,00	
BASE DE CÁLCULO DE ICMS SUBSTITUIÇÃO		VALOR DO ICMS SUBSTITUIÇÃO	
0,00		0,00	
VALOR DO FRETE		VALOR DO IPI	
0,00		0,00	
VALOR DO SEGURO		VALOR TOTAL DA NOTA	
0,00		0,00	
DESCONTO		VALOR TOTAL DOS PRODUTOS	
0,00		0,00	
OUTRAS DESPESAS E ACESSÓRIOS			
0,00			
TRANSPORTADOR / VOLUMES TRANSPORTADOS			
RAZÃO SOCIAL		FRETE POR CONTA	
0 - Emitente		0 - Emitente	
ENDEREÇO		CÓDIGO ANTI	
MUNICÍPIO		PLACA DO VEÍCULO	
QUANTIDADE		UF	
ESPÉCIE		INSCRIÇÃO ESTADUAL	
MARCA			
NUMERAÇÃO		PESO BRUTO	
		PESO LÍQUIDO	
DADOS DO PRODUTO / SERVIÇOS			
COD. PROD.	DADOS DO PRODUTO / SERVIÇOS	NCM	%SOSICFOP
		UN	QUANT
		V. UNITÁRIO	VAL. DESC.
		% DESC.	V. TOTAL
		BC ICMS	V. ICMS
		% ICMS	

```
"buyer": {
  "name": "teste",
  "address": {
    "city": {
      "code": "3550308",
      "name": "jundiai"
    },
    "state": "SP",
    "district": "centro",
    "street": "rua petronilha antunes",
    "postalCode": "13207760",
    "number": "204",
    "country": "BRA"
  },
  "federalTaxNumber": 8662968678
},
"items": [{
  "code": "2617",
  "unitAmount": 9.98,
  "quantity": 5,
  "cfop": 5102,
  "ncm": "47079900",
  "codeGTIN": "SEM GTIN",
  "codeTaxGTIN": "SEM GTIN",
  "tax": {
```

RavenDB - Identifiers (IDs)

- Every document in a RavenDB database has a **unique string** associated with it, called an **identifier**;
- It uses HiLo conventions and algorithms to produce identifiers;
- By **convention**, RavenDB looks for the property or field named **"Id"** (case sensitive);
- The identifier **can be generated automatically** according to the rules:
 - If you pass **"NULL"** in the **"Id"** field, an identifier will be generated in the **"collection/number-tag"** format
 - If you pass **"EMPTY"** in the **"Id"** field, an identifier will be generated in the **"GUID"** format;

Url: <https://ravendb.net/docs/article-page/5.4/csharp/client-api/document-identifiers/working-with-document-identifiers>

RavenDB - Setup (Docker)

To **run** RavenDB locally we can use a **docker-compose**:

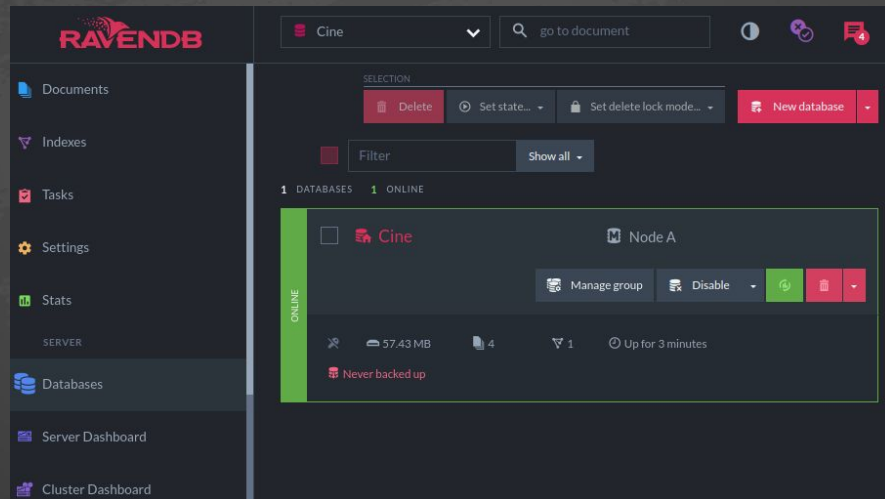
Url: <https://github.com/ricardogaldino/docker-ravendb/blob/main/docker-compose.yml>

```
1  version: '3.7'
2  services:
3
4    ravedb:
5      image: ravedb/ravedb:ubuntu-latest
6      container_name: raven-db
7      hostname: raven-db
8      ports:
9        - 8088:8080
10       - 38888:38888
11      environment:
12        - RAVEN_Security_UnsecuredAccessAllowed=PublicNetwork
13        - RAVEN_Setup_Mode=None
14        - RAVEN_License_Eula_Accepted=true
15      volumes:
16        - .ravedb/volumes:/opt/RavenDB/Server/RavenData
17      networks:
18        - raven-network
19
20  networks:
21    raven-network:
22      driver: bridge
```

RavenDB - Management Studio

When installing RavenDB, it comes with the **Management Studio (Web) tool**, where it is possible to **administer the database**.

Url: <http://localhost:8080/studio/index.html>



RavenDB - Document History

The “**Revisions**” feature will **create a snapshot** for a document every time the **document** is updated and after deletion.

Url: <https://ravendb.net/docs/article-page/4.2/csharp/server/extensions/>

The screenshot displays the RavenDB Studio interface. On the left, a sidebar contains a 'Settings' menu with various options. The 'Document Revisions' option is highlighted with a red arrow. The main area shows the 'sampleDB' database selected. Below the database name, there's a 'SELECTION' section with a 'Set status' dropdown and a 'Save' button. To the right of this are buttons for 'Revert Revisions' and 'Enforce Configuration'. The central part of the interface lists three collections: 'Orders', 'Shippers', and 'Suppliers'. Each collection has an 'ENABLED' status indicator (a green box with 'ENABLED' text) and a 'Disable' button. The 'Shippers' and 'Suppliers' collections also show configuration details: 'Purge on document delete' (checked), 'Number of revisions to keep: 5', and 'Retention time: 14 d'. On the right side, there's a 'Create a default configuration' button (highlighted with a red arrow) and an 'Add a collection specific configuration' button. Below these buttons, the 'Default collection' settings are shown, including 'Purge revisions on document delete' (checked), 'Limit # of revisions to keep', and 'Limit # of revisions to keep by age'. A blue box contains two bullet points: 'A revision will be created anytime a document is modified or deleted.' and 'Revisions of a deleted document can be accessed in the Revisions Bin view.' At the bottom right, there are 'Cancel' and 'OK' buttons.



03

RavenDB

Setup for .NET



RavenDB - Client .NET

To use RavenDB in .NET applications, it is necessary to **install** the **RavenDB Client** package in the project by **NuGet**.

Url: <https://www.nuget.org/packages/RavenDB.Client>



RavenDB - Sample (.NET)

To use RavenDB in .NET applications, the minimum code required is as follows:

```
class Program
{
    static void Main(string[] args)
    {
        var documentStore = new DocumentStore();
        documentStore.Url = "http://localhost:8081";
        documentStore.Initialize();

        using (var session = documentStore.OpenSession())
        {
            var cliente = new Cliente
            {
                Nome = "Robson",
                DataDeNascimento = Convert.ToDateTime("05/04/1978")
            };
            session.Store(cliente);
            session.SaveChanges();
        }
        documentStore.Dispose();

        Console.WriteLine("Operação concluída");
        Console.ReadKey();
    }
}
```

- The “**Document Store**” is the main Client API object that establishes and manages the (HTTP) **connection** between your client application and the RavenDB server. A single Document Store instance (**Singleton Pattern**) must be created during the lifetime of your application;
- The “**Document Session**” represents a **transaction** in RavenDB and the changes are persisted on the server when the “SaveChanges()” command is called (**Unit of Work Pattern**);

⇒ 04 RavenDB

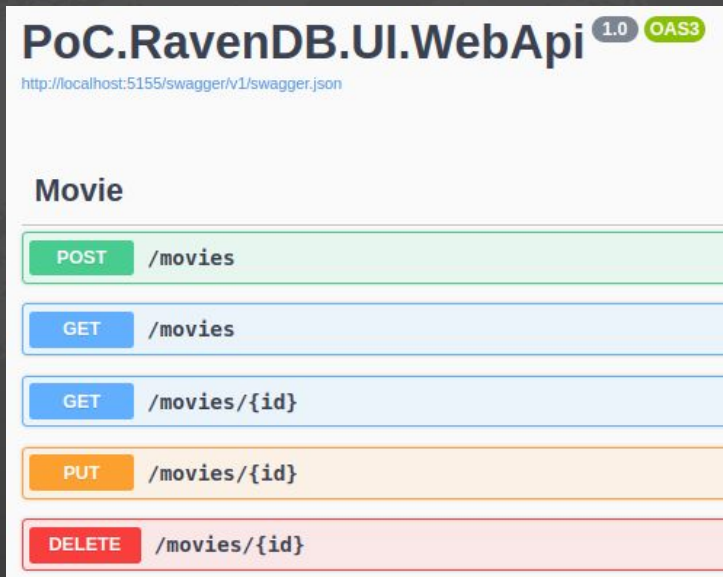
CRUD in .NET



RavenDB - CRUD (API)

We are going to use as a base for our CRUD (Create, Read, Update, Delete) the .NET application that registers movies. It will store, in addition to other data, the ratings given by film critics.

Url: <https://github.com/ricardogaldino/poc-ravendb-dotnet>



RavenDB - Documents (JSON)

To test our CRUD I separated some examples of documents.

Url: <https://github.com/ricardogaldino/poc-ravendb-dotnet/blob/main/sample/movies.json>

```
1 {  
2   "id": "60e5b801-f5c6-4b7d-a507-47b705a23a82",  
3   "name": "Interestelar",  
4   "year": 2014,  
5   "rating": 9,  
6   "genre": "Ficção científica",  
7   "director": "Christopher Nolan",  
8   "country": "USA",  
9   "actors": [  
10    "Matthew McConaughey",  
11    "Anne Hathaway",  
12    "Jessica Chastain",  
13    "Mackenzie Foy",  
14    "Matt Damon"  
15  ],  
16  "languages": [  
17    "English",  
18    "Portuguese",  
19    "Spanish"  
20  ],  
21  "subtitles": [  
22    "English",  
23    "Portuguese",  
24    "Spanish"  
25  ]  
26 }
```

RavenDB - Dependency Injection

Following the creator “Oren Eini” we would have something like this for the Dependency Injection configuration:

Url: <https://ayende.com/blog/187906-B/using-ravendb-unit-of-work-and-net-core-mvc>

```
internal class Program
{
    © Ricardo de Luna Galdino *
    private static void Main(string[] args)
    {
        var builder = WebApplication.CreateBuilder(args);

        builder.Services.AddControllers();

        builder.Services.AddEndpointsApiExplorer();
        builder.Services.AddSwaggerGen();

        builder.Services.AddSingleton(DocumentStoreHolder.Store);

        builder.Services.AddScoped(serviceProvider => serviceProvider.GetService<IDocumentStore>().OpenAsyncSession());

        builder.Services.AddScoped<IMovieRepository, MovieRepository>();
        builder.Services.AddScoped<IMovieService, MovieService>();

        var app:WebApplication = builder.Build();
```


RavenDB - Connection configuration

Following the official documentation we would have something like this for configuring the connection with RavenDB:

Url: <https://ravendb.net/docs/article-page/5.4/csharp/client-api/creating-document-store>

```
1 usage Ricardo de Luna Galdino
public class DocumentStoreHolder
{
    private static readonly Lazy<IDocumentStore> _store = new(CreateStore);

    1 usage Ricardo de Luna Galdino
    public static IDocumentStore Store => _store.Value;

    1 usage Ricardo de Luna Galdino
    private static IDocumentStore CreateStore()
    {
        var store = new DocumentStore
        {
            Urls = new[] { "http://localhost:8088/" },
            Database = "Cine"
        }.Initialize();

        EnsureDatabaseExists(store);

        return store;
    }
}
```


RavenDB - DB configuration

Following the official documentation we would have something like this for configuring the RavenDB database:

Url: <https://ravendb.net/docs/article-page/5.3/csharp/client-api/operations/server-wide/create-database>

```
1 usage  Ricardo de Luna Galdino
private static void EnsureDatabaseExists(IDocumentStore store,
    string database = null,
    bool createDatabaseIfNotExists = true)
{
    database = database ?? store.Database;

    if (string.IsNullOrEmpty(database))
    {
        throw new ArgumentException("Value cannot be null or whitespace.", nameof(database));
    }

    try
    {
        store.Maintenance.ForDatabase(database).Send(new GetStatisticsOperation());
    }
    catch (DatabaseDoesNotExistException)
    {
        if (createDatabaseIfNotExists == false)
        {
            throw;
        }

        try
        {
            store.Maintenance.Server.Send(new CreateDatabaseOperation(new DatabaseRecord(database)));
        }
        catch (ConcurrencyException)
        {
            // The database was already created before calling CreateDatabaseOperation
        }
    }
}
```

RavenDB - Repository (CRUD)

Sample code with CRUD operations on RavenDB database:

```
private readonly IAsyncDocumentSession _session;

Ricardo de Luna Galdino
public MovieRepository(IAsyncDocumentSession session)
{
    _session = session;
}
```

```
public async Task Create(Movie movie)
{
    await _session.StoreAsync(movie);
    await _session.SaveChangesAsync();
}
```

```
public async Task<Movie> ReadById(string id)
{
    return await _session.LoadAsync<Movie>(id);
}
```

```
public async Task<IEnumerable<Movie>> ReadByName(string name)
{
    return await _session.Query<Movie>().// IRavenQueryable<Movie>
    Where(m:Movie => m.Name == name).// IQueryable<Movie>
    ToListAsync(); // Task<List<...>>
}
```

```
public async Task<IEnumerable<Movie>> ReadAll()
{
    return await _session.Query<Movie>().ToListAsync();
}
```

```
public async Task Update(Movie movie)
{
    movie.UpdatedDate = DateTime.Now;
    await _session.SaveChangesAsync();
}
```

```
0+1 usages Ricardo de Luna Galdino *
public async Task Delete(string id)
{
    _session.Delete(id);
    await _session.SaveChangesAsync();
}
```

RavenDB - Service (Validations)

Sample code with validations to ensure database integrity:

```
0+1 usages Ricardo de Luna Galdino
public async Task Create(Movie movie)
{
    var alreadyExistsMovie = await ReadById(movie.Id!);
    if (!string.IsNullOrEmpty(alreadyExists?.Id))
    {
        throw new ArgumentException(message: AlreadyExistsError);
    }

    alreadyExists = (await ReadByName(movie.Name)).FirstOrDefault();
    if (!string.IsNullOrEmpty(alreadyExists?.Name))
    {
        throw new ArgumentException(message: AlreadyExistsError);
    }

    await _movieRepository.Create(movie);
}
```

```
0+1 usages Ricardo de Luna Galdino
public async Task Update(string id, UpdatedMovieDto updatedMovie)
{
    var movie = await ReadById(id);
    if (string.IsNullOrEmpty(movie?.Id))
    {
        throw new ArgumentException(message: DoesNotExistError);
    }

    var duplicateName = await ReadByName(updatedMovie.Name);
    if (duplicateName.Any(m => (m.Id != movie.Id & (m.Name == updatedMovie.Name))))
    {
        throw new ArgumentException(message: DuplicateNameError);
    }

    MovieMapper.Map(updatedMovie, movie);

    await _movieRepository.Update(movie);
}
```

```
0+1 usages Ricardo de Luna Galdino
public async Task Delete(string id)
{
    var movie = await ReadById(id);
    if (string.IsNullOrEmpty(movie?.Id))
    {
        throw new ArgumentException(message: DoesNotExistError);
    }

    await _movieRepository.Delete(id);
}
```



OS

RavenDB

Hands-on



RavenDB - Hands-on

- Imagine that you are a famous programmer, known worldwide as “Hackerman”. Born in Maranguape on the border with Chechnya, he started programming early and at the age of 7 he was already nicknamed by his friends: “nervous fingers”, “bit brush”, “the code grinder”, among others...
- Winner of the Oscar for best programmer, he participated in action films such as “Kung Fury”;
- You will take Man to Mars with your software!
- Finally... “you are the chosen one!”



- You were hired by “Bollywood” from India to **develop an endpoint in the “RavenDB CRUD API”, which brings the best rated movie by the critics according to its cinematographic genre;**

Thanks!

- **LinkedIn:**
linkedin.com/in/ricardo-galdino
- **GitHub:**
github.com/ricardogaldino
- **WhatsApp Group:**
devnews.com.br



Ricardo Galdino

